

Towards an Automated Approach to Apply Idle-time Analysis in Performance Testing

A. Omar Portillo-Dominguez¹, Miao Wang¹, Philip Perry¹,
John Murphy¹, Nick Mitchell², and Peter F. Sweeney²

¹ Lero - The Irish Software Engineering Research Centre, Performance Engineering Laboratory, UCD School of Computer Science and Informatics, University College Dublin, Ireland

`andres.portillo-dominguez@ucdconnect.ie,`
`{philip.perry,miao.wang,j.murphy}@ucd.ie,`

² IBM T.J. Watson Research Center,
Yorktown Heights, New York, USA
`{nickm,pfs}@us.ibm.com`

Abstract. Performance testing in highly distributed environments is a very challenging task. Specifically, the identification of performance issues and the diagnosis of their root causes are time-consuming and complex activities which usually require multiple tools and heavily rely on the expertise of the engineers. In order to simplify the above tasks, hence increasing the productivity and reducing the dependency on expert knowledge, many researchers have been developing tools with built-in expertise knowledge for non-expert users to use. However, variant usability limitations exist in these tools that prevent their efficient usage in performance testing. To address these limitations, this paper presents a lightweight approach to automate the usage of expert tools in performance testing. In this work, we use a tool named Whole-system Idle Time analysis to demonstrate how our research work solves this problem. Our validation involved two case studies using real-life applications, assessing the proposed automation approach in terms of overhead costs and the time savings it brings to the analysis of performance issues. The current results have proven the benefits of the approach by achieving a good decrement in the time invested in performance analysis while generating a low overhead in the tested system.

Keywords: Performance testing, automation, performance analysis, idle-time analysis, distributed environments, multi-tier applications

1 Introduction

It is an accepted fact in the industry that performance is a critical dimension of quality and should be a major concern of any software project. This is especially true at enterprise-level, where performance plays a central role in usability. However it is not uncommon that performance issues occur and materialize into serious problems (i.e. outages on production environments, or even cancellation of projects) in a significant percentage of projects. For example, a 2007 survey applied to information technology executives [1] reported that 50% of them had faced performance problems in at least 20% of their deployed applications.

This situation is partially explained by the pervasive nature of performance, which makes it hard to assess because performance is practically influenced by every single aspect of the design, code, and execution environment of an application. Latest trends in the information technology world (such as Service Oriented Architecture and Cloud Computing) have also augmented the complexity of applications further complicating activities related to performance. Under these conditions, it is not surprising that doing performance testing is complex and time-consuming. A special challenge, documented by multiple authors [2–4], is that current performance tools heavily rely on expert knowledge to understand their output. It is also common that multiple sources are required to diagnose performance problems, especially in highly distributed environments. For instance in Java: thread dumps, garbage collection logs, heap dumps, CPU utilization and JVM memory usage, are a few examples of the information that a tester could need to understand the performance of an application. This situation increases the expertise required to do performance analysis, which is usually held by only a small number of testers within an organization. This could potentially lead to bottlenecks where some activities can only be done by experts, hence impacting the productivity of large testing teams.

To simplify the performance analysis and diagnosis tasks, hence increasing the productivity and reducing the dependency on expert knowledge, many researchers have been developing tools with built-in expertise knowledge for non-expert users [5–7]. However, various limitations exist in these tools that prevent their efficient usage in performance testing. The data collection process usually needs to be controlled manually, which in a highly distributed environment (composed of multiple nodes to monitor and coordinate simultaneously) is very time-consuming and error-prone, specially if the data needs to be processed periodically during the test execution to have an incremental view of the results. The same situation occurs with the outputs, where a tester commonly needs to review multiple reports, one for each monitored node per data processing cycle. Even though these limitations might be manageable in small testing environments, they prevent an effective usage of an expert system in bigger testing environments. For example, considering an application composed of 10 nodes, a 24-hour test run and a monitoring interval of 1 hour to get incremental results. After starting the processes, a tester would need to manually coordinate to stop the data collections, generate the outputs of the expert system and then start the data collections again. These steps conducted for the 10 nodes in cycles of 1 hour, which throws a total of 240 cycles. Additionally, these steps would need to be carried out as fast as possible to minimize the time gaps between the end of a cycle and the start of the next. Lastly, the tester would have to review the 10 different reports she would get every hour and compare them with the previous ones to evaluate if there are any performance issues. As it can be inferred from this example, the costs of using an expert system in a highly distributed environment would outweigh the benefits, as the time and effort required to synchronize its execution and the analysis of its outputs would be very effort-intensive and error-prone. These limitations could also lead testers to take major assumptions

such as focusing the analysis on a single randomly selected node, assuming it is representative of the whole system.

As a highly distributed environment is precisely the scenario where the analysis capabilities of an expert system would be more valuable, these usage limitations must be addressed.

In addition to the previous challenges, the overhead generated by any technique should be low to minimize the impacts it has in the tested environment (i.e. inaccurate results, abnormal behaviours), otherwise the technique would not be suitable for performance testing. For example, instrumentation³ is currently a common approach used in performance analysis to gather input data [8–11]. However, it has the downside of obscuring the performance of the instrumented applications, hence compromising the results of the performance testing. Similarly, if a tool requires heavy human effort to be used, this might limit the applicability of that tool. On the contrary, automation could play a key role to encourage the adoption of a technique. As documented by the authors in [12], this strategy has proven successful in performance testing activities.

Finally, to ensure that our research work is helpful for solving real-life problems in the software industry, we have been working with our industrial partner, the IBM System Verification Test (SVT), to understand the challenges that they experience in their day-to-day testing activities. The received feedback confirms that there is a real need to simplify the usage of expert tools so that testers can carry out analysis tasks in less time.

This paper proposes a lightweight automation approach that addresses the common usage limitations of an expert system in performance testing. Furthermore, during our research development work we have successfully applied our approach to the IBM Whole-system Analysis of Idle Time tool (WAIT)⁴. This freely available tool is a lightweight expert system that helps to identify the main performance inhibitors that exist on Java-based systems. Our work was validated through two case studies using real-world applications. The first case study concentrated in evaluating that the overhead introduced by our approach was low. The second assessed the productivity benefits that the proposed approach brings to the performance testing process. The current results have provided evidence about the benefits of this approach: Using the automated approach drastically reduced the effort required by a tester to use and analyze the outputs of the selected expert system (WAIT). This usage simplification translated into a quicker identification of performance issues, including the pinpointing of the responsible classes and method calls. Moreover the overhead in the monitored system was low (between 0% to 3% when using a common industry *Sampling Interval*).

The main contributions of this paper are:

1. A novel lightweight approach to automate the usage of expert systems in performance testing.

³ [http://msdn.microsoft.com/en-us/library/aa983649\(VS.71\).aspx](http://msdn.microsoft.com/en-us/library/aa983649(VS.71).aspx)

⁴ <http://wait.ibm.com>

2. A practical validation of the approach consisting of its implementation around the WAIT tool and two case studies using real-life applications.
3. An analysis of the overhead this approach has in the monitored environment to demonstrate it has a low impact.

The rest of this paper is structured as follows: Section 2 discusses the background. Section 3 explains the proposed approach, while Section 4 explores the experimental evaluation and results. Section 5 shows the related work. Finally Section 6 presents the conclusions and future work.

2 Background

Idle-time analysis is a methodology that pursues to explain the root causes that lead to under-utilized resources. This approach, proposed in [5], is based on the observed behavior that performance problems in multi-tier applications usually manifest as idle time indicating a lack of motion. WAIT is an expert system that implements the idle-time analysis and identifies the main performance inhibitors that exist on a system. It has proven successful in simplifying the detection of performance issues and their root causes in Java environments [5, 13]. Moreover WAIT is based on non-intrusive sampling mechanisms available at Operating System level (i.e. “ps” and “vmstat” commands in a Unix/Linux environment) and the Java Virtual Machine (JVM), in the form of *Javacores*⁵ (diagnostic feature of Java to get a quick snapshot of the JVM state, offering information such as threads, locks and memory). The fact that WAIT uses standard data sources makes it non-disruptive, as no special flags or instrumentation are required to use it. Moreover WAIT requires infrequent samples to perform its diagnosis, so it also has low overhead.

From a usability perspective, WAIT only requires a few steps: A user needs to collect as much data as desired, upload it to a web page and obtain a report with the findings. This process can be repeated multiple times to monitor a system through time. Internally, WAIT uses an engine built on top of a set of expert rules that perform the analysis. Fig. 1 shows an example of a WAIT Report. The top area summarizes the usage of resources (i.e. CPU, thread pools and Java Heap) and the number and type of threads per sample. The bottom section shows all the performance inhibitors that have been identified, ranked by frequency and impact. Moreover each category of problem is indicated with a different color. For example, in Fig. 1 the top issue appeared in 53% of the samples and affected 7.6 threads on average. The affected resource was the network (highlighted in yellow) and was caused by database readings.

Given its strengths, WAIT is a promising candidate to reduce the expert knowledge and time required to do performance analysis. However, it has some usability limitations that exemplify the current challenges to use effectively an expert system in the performance testing domain: The effort required to do manual data collection to feed WAIT and the number of WAIT reports a tester

⁵ <http://www-01.ibm.com/support/docview.wss?uid=swg27017906&aid=1>

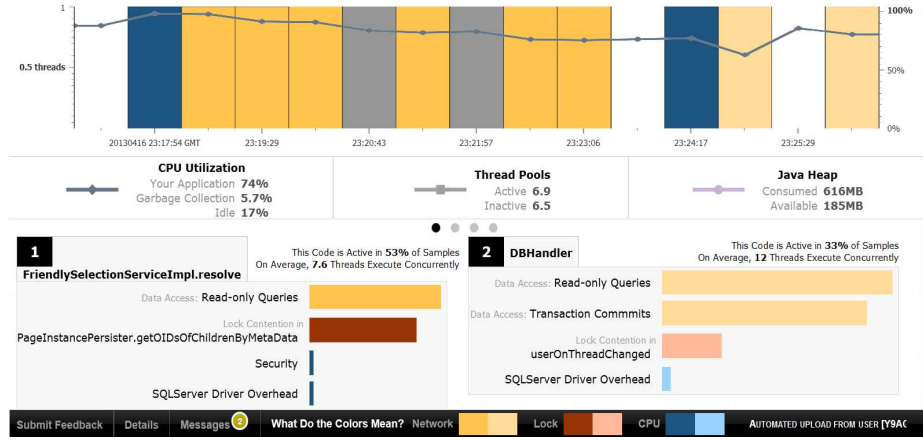


Fig. 1. Example of WAIT Report

would need to review are practically linear with respect of the number of nodes in the application and the frequency with which the data is refreshed. These costs make the usage of WAIT very effort-intensive and error-prone in a highly distributed environment, making it an ideal candidate to apply our proposed approach.

3 Proposed Approach and Implementation

This section presents the proposed approach, the architecture that supports it and the prototype that has been developed.

3.1 Proposed Approach

The objective of this work is to address the common usability limitations of expert systems to be applied effectively in performance testing. To achieve this, our approach proposes the automation of the manual processes involved in the usage of an expert system: It will execute jointly to a performance test, periodically collecting the required samples, then getting them incrementally processed in order to get a consolidated output.

This approach is depicted in the Procedure Proposed Approach and requires a few inputs: The list of nodes to be monitored; a *Sampling Interval* to control how often the samples will be collected; a *Time Threshold* to define the maximum time between uploads; a *Hard Disk Threshold* to define the maximum storage quota for collected data (to prevent its uncontrolled growth); and a *Backup* flag that indicates if the collected data should be backed up before any cleaning occurs.

The process starts by getting a new *RunId*, value which will uniquely identify the test run and its reports. This value is propagated to all the nodes. On each node, the *Hard Disk Usage* and the *Next Time Threshold* are initialized. Then each node starts in parallel the following loop until the performance test finishes: A new set of data samples is collected. After the collection finishes,

Procedure Proposed Approach

Input: A set AN of n application nodes, Sampling Interval $Sint$, Time Threshold $tThres$, Hard Disk Threshold $hdThres$, Backup Flag $bFlag$. If $bFlag = true$, Backup destination path $bPath$.

Output: Consolidated Outputs

// Initialization

- 1 obtain $runId \leftarrow \text{new RunId}$
- 2 share $runId$ with all nodes
- 3 **foreach** $node$ in AN (*in parallel*) **do**
- 4 $hdUsage \leftarrow 0$
- 5 $nextTimeThreshold \leftarrow \text{current time from the Operating System} + tThres$
- 6 **// Main Process**
- 7 **while** *performance test is executing* **do**
- 8 **// Data gathering**
- 9 collect new set of data samples
- 10 **// Update performance results**
- 11 $currentTime \leftarrow \text{current time from the Operating System}$
- 12 $hdCurrentUsage \leftarrow \text{calculate Hard Disk space of collected data}$
- 13 **if** $currentTime \geq nextTimeThreshold$ **or** $hdCurrentUsage \geq hdThres$ **then**
- 14 send locally collected data indicating it as part of test $runId$
- 15 **if** $bFlag = true$ **then**
- 16 copy locally collected data to $bPath$ indicating it as part of test $runId$
- 17 delete locally collected data
- 18 retrieve consolidated outputs
- 19 $nextTimeThreshold \leftarrow \text{current time from the Operating System} + tThres$
- 20 wait $Sint$ before performing next iteration of the process
- 21 **// Closure**
- 22 send remnant locally collected data indicating it as part of test $runId$
- 23 **if** $bFlag = true$ **then**
- 24 copy remnant locally collected data to $bPath$ indicating it as part of test $runId$
- 25 delete remnant locally collected data
- 26 retrieve final consolidated outputs

it is assessed if any of the two thresholds has been reached (either the *Hard Disk Usage* has exceeded the *Hard Disk Threshold* or the *Next Time Threshold* has been reached). If any of the conditions has occurred, a new upload round occurs where the data is share with the expert system (labeling the data with the *RunId* so that information from different nodes can be identified as part of the same test run). If a *Backup* was enabled, the data is copied to the backup destination before it is deleted. Then the updated outputs of the expert system are retrieved and the *Next Time Threshold* is calculated. Finally, the logic awaits the configured *Sampling Interval* before a new iteration starts.

Once the performance test finishes, any remnant collected data is send (and backup if configured) so that this information is also considered by the expert system. Then this data is also cleared and the final consolidated outputs of the expert system are obtained.

3.2 Architecture

The previously described approach was complemented with the architecture presented in Fig. 3. It is composed of two main components: The *Control Agent* is responsible of interacting with the Load Testing tool to know when the performance test starts and ends. It is also responsible of getting the runId and propagate it to all the nodes. The second component is the *Node Agent* which is responsible of the collection, upload, backup and cleanup steps in each application node.

[PENDING]

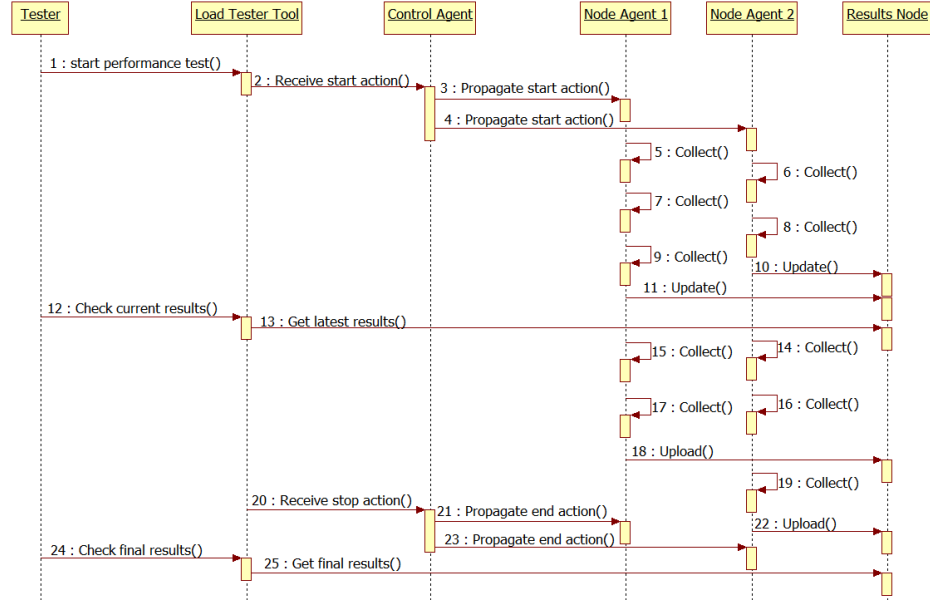


Fig. 2. Sequence Diagram of automated WAIT

4 Experimental Evaluation

In this section the developed prototype and the experimental setup are presented. Then the experiments are described and their results discussed.

4.1 Prototype

Based on the proposed approach, a prototype has been developed in conjunction with our industrial partner IBM. The *Control Agent* was implemented as an

Eclipse Plugin for the Rational Performance Tester (RPT) ⁶, which is a load testing tool commonly used in the industry; the *Node Agent* was implemented as a Java Web Application, and WAIT was the selected expert system due to its analysis capabilities to diagnose performance issues and their root causes.

Once the agents are installed, WAIT can now be configured as any other resource in RPT as shown in Fig. 4. Similarly, during a performance test WAIT can be monitored as any other resource in the *Performance Report* of RPT under the *Resource View* depicted in Fig. 5. Finally, the consolidated WAIT report is also accessible within RPT, so a tester does not need to leave RPT during the whole duration of the performance test. This is shown in Fig. 6.

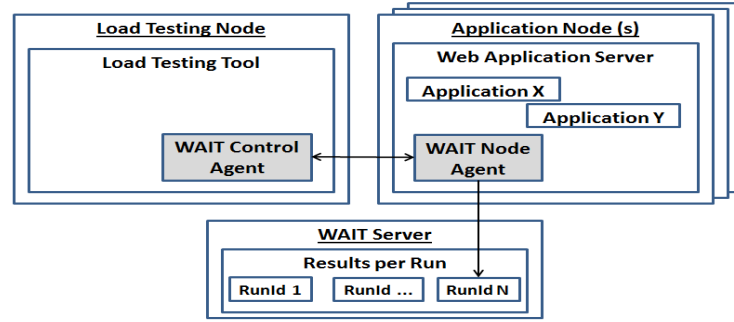


Fig. 3. High-level Architecture of automated WAIT

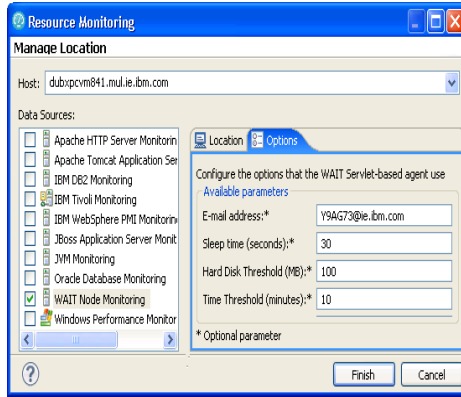


Fig. 4. WAIT configuration in RPT

The screenshot shows the 'Performance Report' window in RPT, displaying resources for the time range of 2000 users. The table lists various metrics for the resource 'dubxpcvm841'.

Resources (Displaying For Time Range: 2000 Users)	dubxpcvm841
WAIT -- Monitored PIDs (Average for Run)	1
WAIT -- Monitored PIDs (Minimum for Run)	1
WAIT -- Monitored PIDs (Maximum for Run)	1
WAIT -- Collection Iterations (Average for Run)	10
WAIT -- Collection Iterations (Minimum for Run)	10
WAIT -- Collection Iterations (Maximum for Run)	10
WAIT -- Successful collections (Average for Run)	9
WAIT -- Successful collections (Minimum for Run)	9
WAIT -- Successful collections (Maximum for Run)	9

Fig. 5. WAIT monitored in RPT

4.2 Experimental Set-up

Two experiments were performed. The first one aimed to evaluate if the overhead introduced by the proposed approach was low to prevent compromise the results of a performance test execution. Meanwhile, the second experiment pursued to assess the productivity benefits that a tester can gain by using the proposed approach.

⁶ <http://www-03.ibm.com/software/products/us/en/performance>

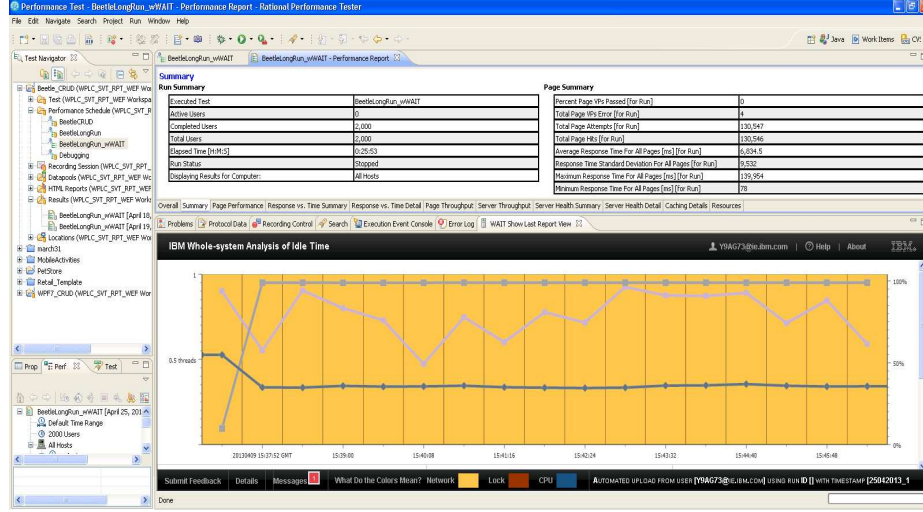


Fig. 6. WAIT Report accessible in RPT

Moreover two environment configurations were used: One was composed of one RPT node, one application node and one *WAIT Server* node; the other was composed of one RPT node, one load balancer node, two application nodes and one *WAIT Server* node. All connected by a 10-Gbit LAN. These set-ups are shown in Fig. 7.




Fig. 7. Environment Configurations

The RPT node ran over Windows XP with an Intel Xeon CPU at 2.67 Ghz and 3GB of RAM using RPT 8.2.1.3. The *WAIT Server* was run over Red Hat Enterprise Linux Server 5.9, with an Intel Xeon CPU at 2.66 GHz and 2GB of RAM using Apache Web Server 2.2.3. Each application node was a 64-bit Windows Server 2008, with an Intel Xeon E7-8860 CPU at 2.26 GHz and 4GB of RAM running Java 1.6.0 IBM J9 VM (build 2.6). Finally, the load balancer node had the same characteristics of the *WAIT Server* node.

4.3 Experiment #1: Overhead Evaluation

Its objective was to validate that the proposed approach had low overhead and involved the assessment of four metrics: Throughput (hits per second), response time (milliseconds), CPU (%) and memory (MB) utilization. All metrics were collected through RPT. Furthermore, 2 real-world applications were used: iBatis

JPetStore 4.0 ⁷ which is an upgraded version of Sun’s Pet Store, an e-commerce shopping cart application. It ran over an Apache Tomcat 6.0.35 Application Server. The other application was IBM WebSphere Portal Server 8.0.1 ⁸, a leader solution in the enterprise portal market [14]. It ran over an IBM WebSphere Application Server 8.0.0.5.

Firstly, the overhead of the approach was measured in a single-node environment. For each application, 3 combinations of WAIT were evaluated: The application alone to get a baseline; the application with manual WAIT data collection; and the application with automated WAIT. For each combination using WAIT, the *Sampling Intervals* were configured to 480 seconds (suggested by IBM SVT) and 30 seconds (minimum value recommended for WAIT). The remaining configuration parameters were suggested by IBM SVT to reflect real-world conditions: A workload of 2,000 concurrent users; a duration of 1 hour; a *Hard Disk Threshold* of 100MB; and a *Time Threshold* of 10 minutes. Finally, each combination was repeated 3 times.

For JPetStore, each test run produced around 500,000 transactions. The results presented in Table 1 showed that using WAIT with a *Sampling Interval* of 480 seconds had practically no impact in terms of response time and throughput. Furthermore the difference in resource consumption between the different modalities of WAIT was around 1%. This difference was mostly related to the presence of the *WAIT Node Agent* because the uploaded data was very small in this case (around 200KB every 10 minutes). When WAIT used a *Sampling Interval* of 30 seconds, the impacts in response time and throughput appeared. Considering the throughput was similar between the WAIT modalities, the impact was caused by the *Javacore* generation (step shared between the modalities). In average, the generation of a *Javacore* took around 1 second. Even though this cost was insignificant in the higher *Sampling Interval*, with 30 seconds the impact was visible. On the contrary, the difference in response times (2.8%, around 53 milliseconds) was caused by the upload and backup processes (around 4MB of data every 10 minutes), as the cost of the *WAIT Node Agent* presence had been previously measured. In terms of resource consumption, the differences between the WAIT modalities remained within 1%.

Table 1. JPetStore - Overhead Results

WAIT Modality	Avg Response Time (ms)	Max Response Time (ms)	Avg Throughput (hps)	Avg CPU Usage (%)	Avg Memory Usage (MB)
None (<i>Baseline</i>)	1889.6	44704.0	158.8	36.9	1429
Manual, 480s	0.0%	0.0%	0.0%	1.1%	3.0%
Automated, 480s	0.0%	0.0%	0.0%	2.0%	3.7%
Manual, 30s	1.6%	0.4%	-4.0%	1.47%	4.1%
Automated, 30s	4.4%	0.5%	-3.1%	2.53%	4.4%

⁷ <http://sourceforge.net/projects/ibatisjpetstore/>

⁸ <http://www-03.ibm.com/software/products/us/en/portalserver>

For Portal, each test run produced around 400,000 transactions. Even though the results presented in Table 2 showed similar trends in terms of having lower overheads using the *Sampling Interval* of 480 seconds, a few key differences were identified: First, the impacts in response time and throughput were visible since the *Sampling Interval* of 480 seconds. Besides, the differences between *Sampling Intervals* were bigger. As the experiment conditions were the same, it was initially assumed that these differences were related to the dissimilar functionality of the tested applications. This was confirmed after analyzing the *Javacores* generated by Portal, which allowed to quantify the differences in behavior of Portal: The average size of a *Javacore* was 5.5MB (450% bigger than JPetStore's), its average generation time was 2 sec (100% bigger than JPetStore's), with a maximum generation time of 3 sec (100% bigger than JPetStore's).

Table 2. Portal - Overhead Results

WAIT Modality	Avg Response Time (ms)	Max Response Time (ms)	Avg Throughput (hps)	Avg CPU Usage (%)	Avg Memory Usage (MB)
None (<i>Baseline</i>)	4704.75	40435.50	98.05	76.73	3171.20
Manual, 480s	0.7%	0.6%	-0.1%	1.13%	2.2%
Automated, 480s	3.4%	1.0%	-2.8%	0.63%	4.1%
Manual, 30s	14.9%	5.4%	-5.7%	2.97%	5.3%
Automated, 30s	16.8%	9.1%	-5.6%	2.23%	6.0%

Due to the small differences among the runs and the variations (presumable environmental) that were experienced during the experiments, a Paired t-Test⁹, was done (using a significant level of $p < 0.1$) to evaluate if the differences in response time and throughput were statistically significant. The results presented in Table 3 showed that for JPetStore the only significant differences existed in the average response time and the average throughput (automated WAIT) when using a *Sampling Interval* of 30 seconds. Similar results were obtained from Portal. This analysis reinforced the conclusion that the overhead was low and the observation that the *Sampling Interval* of 480 seconds was preferable.

A second test was done to validate that the overhead of the proposed approach remained low when used in a multi-node environment over a longer test run. This test used JPetStore and the automated WAIT with a *Sampling Interval* of 480 seconds. The rest of the set-up was identical to the previous tests except the workload which was doubled to compensate the additional application node and the test duration which was increased to 24 hours. Even though the results were slightly different than the single-node run, they proved that the solution was reliable, as using the automated approach had minimal impact in terms of response time (0.5% in the average and 0.2% in the max) and throughput (1.4%). Moreover the consumption of resources behaved similar to the single-node test (an increment of 0.85% in CPU and 2.3% in Memory).

⁹ <http://www.aspfree.com/c/a/braindump/comparing-data-sets-using-statistical-analysis-in-excel/>

Table 3. Paired t-Test Results

Application	WAIT Modality	Avg Response Time (ms)	Max Response Time (ms)	Avg Throughput (hps)
JPetStore	Manual, 480s	0.470	0.143	0.206
JPetStore	Automated, 480s	0.342	0.297	0.472
JPetStore	Manual, 30s	0.089	0.241	0.154
JPetStore	Automated, 30s	0.019	0.334	0.078
Portal	Manual, 480s	0.140	0.263	0.496
Portal	Automated, 480s	0.040	0.189	0.131
Portal	Manual, 30s	0.001	0.158	0.167
Portal	Automated, 30s	0.013	0.105	0.072

The performed paired t-Test also indicated that the differences in response time and throughput between the test runs were not statistically significant.

In conclusion, the results of this first experiment showed that the overhead caused by the automated approach remained low. This answered our research question Q3 positively (Can the overhead, introduced by the proposed approach, be low enough such that it does not compromise the results during a performance test execution?), but with a side note: Due to the impact that the *Sampling Interval* and the application behavior could have in the overhead, it is important to consider these factors when using WAIT. In our case, a *Sampling Interval* of 480 seconds proved efficient in terms of overhead for the two tested applications.

4.4 Experiment #2: Assessment of productivity benefits

Here the objective was to assess the benefits our approach brings to a performance tester. First, the source code of JPetStore was modified and 3 common performance issues were injected: A lock contention bug, composed by a very heavy calculation within a synchronized block of code; a I/O latency bug, composed by a very expensive file reading method; and a deadlock bug, an implementation of the classic “friends bowing” deadlock example¹⁰. Then the automated WAIT was used to monitor this application to assess how well it was able to identify the injected bugs and estimate the corresponding time savings in performance analysis. All set-up parameters were identical to the multi-node test previously performed with exception of the duration which was reduced to 1-hour. Due to space constraints, only the most relevant sections of the WAIT report are presented.

Surprisingly the 1st ranked issue was none of the injected bugs but a method named “McastServiceImpl.receive” which appeared in practically all the samples. Further analysis determined this method call was benign and related to the clustering functionality of Tomcat. The 2nd ranked issue was the lock contention. A relevant point to highlight is that both issues were detected since the early versions of the report. Based on their high frequency (above 96% of the samples), this information could have led a tester to pass this information to the development team so that the diagnosis could start far ahead of the test

¹⁰ <http://docs.oracle.com/javase/tutorial/essential/concurrency/deadlock.html>

completion. The final report reinforced the presence of these issues by offering similar ranking. Fig. 8.a shows the results of the early report, while 8.b shows the results of the final report.

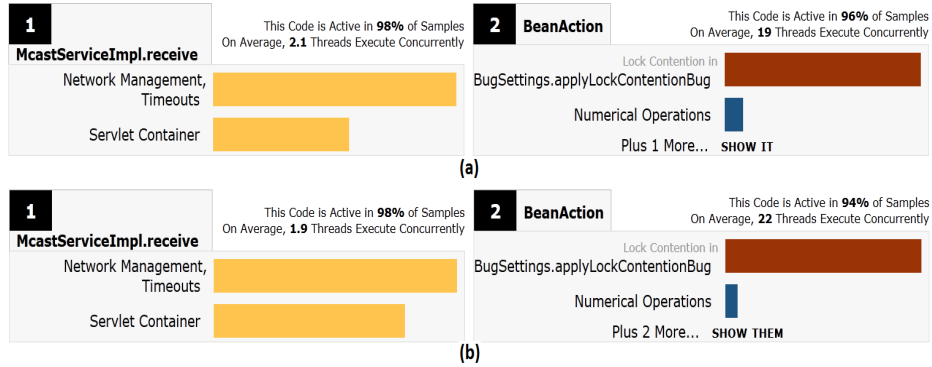


Fig. 8. Top detected performance issues in modified JPetStore application

After identifying an issue, a tester can see more details, including the type of problem, involved class, method and method line. Fig. 9 shows the information of our Lock Contention bug, which was located in the LockContentionBug class, the method generateBug and the line 20. When comparing this information with the actual code, one can see that is precisely the line where the bug was injected (taking a class lock before doing a very CPU intensive logic).

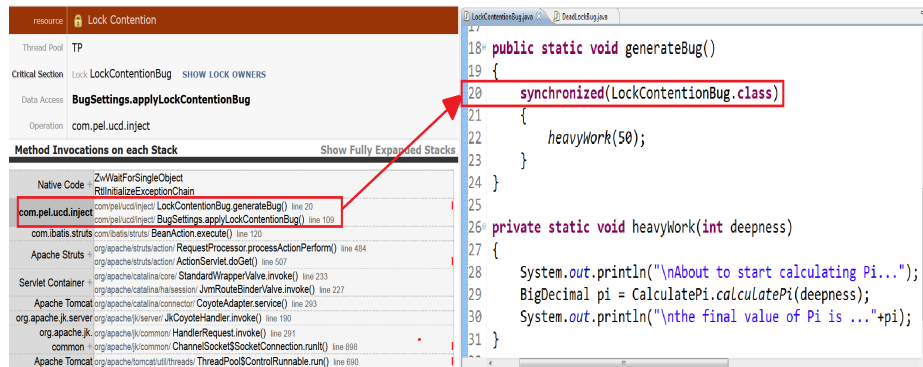


Fig. 9. Lock contention issue in the WAIT report and the actual source code

In 3rd place the report showed a symptom of the lock contention issue (the issues were correlated by comparing their detail information, which pinpointed to the same class/method), suggesting this was a major issue. The I/O latency bug was identified in 4th place. Fig. 10 shows the details of these issues.

The deadlock issue did not appear in this test run, somehow prevented by the lock contention bug which had a major impact that planned. As in any common test phase, the identified bugs were “fixed” and a new run was done to review if any remaining performance issues existed. Not surprisingly, the deadlock bug

CatalogService	SystemLogHandler.println
resource: Lock Contention	resource: Disk
Thread Pool: TP	Thread Pool: com.pel.ucd.inject.HeavyFileReadingBug\$1
Critical Section: Lock LockContentionBug SHOW LOCK OWNERS	Data Access: (to Filesystem)
Data Access: ProductSqlMapDao.getProduct	Operation: Filesystem Data Access
Operation: com.pel.ucd.inject	
Method Invocations on each Stack	Method Invocations on each Stack
Native Code: ZwWaitForSingleObject RtlInitializeExceptionChain	Native Code: ZwWaitForSingleObject jtiAcquireVMAccess
com.pel.ucd.inject: LockContentionBug.generateBug() line 20	Filesystem Data Access: java.io.FileOutputStream.writeBytes() native method
com.pel.ucd.inject: BugSettings.applyLockContentionBug() line 109	Access: java.io.FileOutputStream.write() line 262
com.ibatis.petstore.persistence: ProductSqlMapDao.getProduct() line 30	Java Standard Library: java.io.BufferedOutputStream.flushBuffer() line 65
com.ibatis.dao.engine.impl: DaoProxy.invoke() line 61	Parsing and Formatting: java.io.PrintStream.write() line 445
com.ibatis.dao.engine.proxy: \$Proxy4.getProduct() bytecode 18	Java Standard Library: sun.nio.cs.StreamEncoder\$CharsetSE.writeBytes() line 343
com.ibatis.petstore.service: CatalogService.getProduct() line 60	Parsing and Formatting: sun.nio.cs.StreamEncoder.flushBuffer() line 126
com.ibatis.petstore.presentation: CatalogBean.viewProduct() line 165	Output: java.io.OutputStreamWriter.flushBuffer() line 176
com.ibatis.struts: BeanAction.execute() line 143	Formatting: java.io.PrintStream.println() line 770
	com.ibm.jvm.io: ConsolePrintStream.println() line 299
	Apache Tomcat: org.apache.tomcat.util.log.SystemLogHandler.println() line 238
	com.pel.ucd.inject: HeavyFileReadingBug.read() line 104
	com.pel.ucd.inject: HeavyFileReadingBug\$1.run() line 44

Fig. 10. Details of issues ranked 3rd and 4th

appeared. Fig. 11 shows the information of our Deadlock bug, which was located in the line 30 of the DeadLockBug class. This is precisely the line where the bug was injected (as the deadlock occurs when the friends bow back to each other).

unknown	This Code is Active in 72% of Samples On Average, 41 Threads Execute Concurrently
Thread Pool: com.pel.ucd.inject.DeadLockBug\$1	
Operation: com.pel.ucd.inject	
Method Invocations on each Stack	
Native Code: ZwWaitForSingleObject RtlInitializeExceptionChain	
com.pel.ucd.inject: DeadLockBug\$Friend.bow() line 30	
com.pel.ucd.inject: DeadLockBug\$1.run() line 55	

<pre> 20 21 public String getName() 22 { 23 return this.name; 24 } 25 26 public synchronized void bow(Friend bower) 27 { 28 System.out.format("%s: %s" + " has bowed to me!\n", 29 this.name, bower.getName()); 30 bower.bowBack(this); 31 } 32 33 public synchronized void bowBack(Friend bower) 34 { </pre>

Fig. 11. Deadlock issue in the WAIT report and the actual source code

This experiment was considered successful as all injected bugs were identified, including the involved classes and methods. In terms of time, two main savings were documented. First, the automated approach practically reduced the effort of collecting data for WAIT to zero. After a one-time installation which took no more than 15 minutes for all nodes, the only additional time required to use the automate approach were a few seconds required to modify its configuration (i.e. to change the *Sampling Interval*). The second time saving was in the analysis of the WAIT reports. Previously, a tester would have ended with multiple reports. Now a tester only needs to monitor a single report which is refreshed periodically.

Overcoming the usability constraints of WAIT also allow a tester to exploit WAIT's expert knowledge capabilities. Even though it might be hard to define an average time spent identifying performance issues, a conservative estimate (i.e. 2 hours per bug) could help to quantify WAIT's savings. In our case study, instead of spending an estimate of 6 hours analyzing the issues, it was possible to identify them and their root causes in a matter of minutes with the information provided by the WAIT report. As seen in the experiment, additional time can

also be saved if relevant issues are reported to developers in parallel to the test execution. This is especially valuable in long-term runs, common in performance testing and which usually last several days.

In conclusion, the results of this second experiment answered our research question Q2 (What are the productivity benefits that a tester can gain by using the proposed approach?), showing how a tester’s productivity improves by using an automated WAIT.

To summarize the experimental results, they were satisfactory because it was possible to achieve the goal of automating the execution of WAIT, while keeping the overhead low (in the range of 0% to 3% using a *Sample Interval* of 480 seconds). Additionally the automated approach brought several time savings to a tester: After a quick installation (around 5 minutes per node), the time required to use an automated WAIT was minimal. Also a tester now only needs to monitor a single WAIT report, which offers a consolidated view of the results. A direct result of these savings is the reduction in effort and expert knowledge required by tester to identify performance issues, hence improving the productivity.

4.5 Threats to Validity

Like any empirical work, there are some threats to the validity of these experiments. First the possible environmental noise that could affect the test environments because they are not isolated. To mitigate this, multiple runs were executed for each identified combination. Another threat was the selection of the tested applications. Despite being real-world applications, their limited number implies that not all types of applications have been tested and wider experiments are needed to get more general conclusions. However, there is no reason to believe that the approach is not applicable to other environments.

5 Related Work

The idea of applying automation in the performance testing domain is not new. However, most of the research has focused on automating the generation of load test suites [15–22]. Regarding performance analysis, a high percentage of the proposed techniques require instrumentation. For example, the authors in [8] instrument the source code to mine the sequences of call graphs to infer any relevant error patterns. A similar case occurs with the work presented in [9, 10] which rely on instrumentation to dynamically infer invariants and detect programming errors; or the approach proposed by [11] which uses instrumentation to capture execution paths to determine the distributions of normal paths and look for any significant deviations to detect errors. In all these cases, instrumentation would obscure the performance of an application during performance testing hence discouraging their usage. On the contrary, our proposed approach does not require any instrumentation.

The closest work to ours is [23]. Here the authors present a non-intrusive approach which automatically analyzes the execution logs of a load test to identify performance problems. As this approach only relies on load testing results, it can not determine root causes. Another approach related to our work is presented in

[24] which aims to offer information about the causes behind the issues. However it only limits to provide the subsystem responsible of the performance deviation. On the contrary, our approach allows the applicability of the idle-time analysis in the performance testing domain, through automation means, which allows to identify the classes and methods responsible of the performance issues. Moreover the techniques presented in [23, 24] require information from previous runs as baseline to do their analysis, information which might not always be available.

6 Conclusions and Future Work

The identification of performance problems and the diagnosis of their root causes in highly distributed environments are complex and time-consuming tasks, which tend to rely on the expertise of the engineers involved. The objective of this work was to reduce the expert knowledge and effort required to identify performance issues and their root causes. To achieve this goal a novel automation approach was proposed that automatically gathered information about resource idle time across a distributed computing infrastructure. This technique was implemented in a prototype tool and the tools overhead impact was assessed against the performance of a non-monitored system.

The results showed that for a *Sampling Interval* of 480 seconds, the fully automated system caused a negligible degradation to the system response time and throughput for the JPetstore application. For the Portal application it caused a 3.4% increase in the average response time and a 2.8% reduction in throughput. The current results have also provided evidence of the time savings in performance testing activities that can be achieved by applying the proposed approach to existing expert tools.

In our case, the effort required to use WAIT depended on the number of nodes that composed the monitored environment and the frequency with which result updates were required. These dependencies made WAIT usage very time-consuming and error-prone in highly distributed environments. After applying the proposed automation approach, the effort required to use WAIT was reduced to seconds, regardless of the number of nodes or the desired frequency of result updates. This usage simplification provoked that the effort required by a tester to identify defects and their root causes was minimized: In our case study with a modified JPetStore application, the three injected defects and their root causes were identified in a matter of minutes. In contrast, a similar diagnosis analysis might have taken around 6 hours if performed manually (considering a conservative estimate of 2 hours per bug).

Thus, the approach was shown to have a low overhead in these test cases and was shown to provide an easily accessible summary of the system performance presented in a single report. This tool therefore has the possibility of reducing the time required to analyse performance test data and thereby reduce the effort and costs associated with software system performance testing.

Future work will concentrate on assessing the approach and its benefits through broader case studies with our industrial partner IBM with a special interest in understanding the trade-off between the *Sampling Interval* and the

nature of the applications (represented by their *Javacores*). It will also be investigated how best to exploit the functional information that can now be obtained from a tested environment (i.e. workload, throughput or transactions) to improve the qualitative and quantitative capabilities of the idle-time analysis to identify more types of performance problems.

Acknowledgments

We would like to thanks Amarendra Darisa, from IBM SVT, as his experience in performance testing helped us through the scope definition and validation. This work was supported, in part, by Science Foundation Ireland grant 10/CE/I1855 to Lero - the Irish Software Engineering Research Centre (www.lero.ie).

References

1. Compuware: Applied Performance Management Survey. (2007)
2. Woodside, M., Franks, G., Petriu, D.C.: The Future of Software Performance Engineering. Future of Software Engineering (FOSE '07) (May 2007) 171–187
3. Parsons, T., Murphy, J.: Detecting Performance Antipatterns in Component Based Enterprise Systems. In: 2nd International Middleware Doctoral Symposium. Volume 7. (2008) 55–90
4. Angelopoulos, V., Parsons, T., Murphy, J., O'Sullivan, P.: GcLite: An Expert Tool for Analyzing Garbage Collection Behavior. 2012 IEEE 36th Annual Computer Software and Applications Conference Workshops (July 2012) 493–502
5. Altman, E., Arnold, M., Fink, S., Mitchell, N.: Performance analysis of idle programs. ACM SIGPLAN Notices **45**(10) (October 2010) 739
6. Ammons, G., Choi, J.d., Gupta, M., Swamy, N.: Finding and Removing Performance Bottlenecks in Large Systems. In: ECOOP 2004 Object-Oriented Programming. (2004)
7. Angelopoulos, V., Parsons, T., Murphy, J., O'Sullivan, P.: GcLite: An Expert Tool for Analyzing Garbage Collection Behavior. IEEE Annual Computer Software and Applications Conference Workshops (July 2012) 493–502
8. J. Yang, D. Evans, D.T.M.: Perracotta: mining temporal api rules from imperfect traces. International conference on Software engineering (2008)
9. S. Hangal, M.: Tracking down software bugs using automatic anomaly detection. International Conference on Software Engineering (2002)
10. C. Csallner, Y.: Dsd-crasher: a hybrid analysis tool for bug finding. International symposium on Software testing and analysis (2006)
11. M. Y. Chen, A. Accardi, E.J.A.E.: Path-based failure and evolution management. Symposium on Networked Systems Design and Implementation (2004)
12. Shahamiri, S.R., Kadir, W.M.N.W., M.: A Comparative Study on Automated Software Test Oracle Methods. ICSEA (2009)
13. Wu, Haishan, Asser N. Tantawi, T.: A Self-Optimizing Workload Management Solution for Cloud Applications. (2012)
14. Gootzit, David, Gene Phifer, R.: Magic Quadrant for Horizontal Portal Products. Technical report, Gartner Inc. (2008)
15. Chen, S., Moreland, D., N.Z.: Yet Another Performance Testing Framework. Australian Conference on Software Engineering (ASWEC) (2008)

16. Albert, Elvira, Miguel Gmez-Zamalloa, J.: Resource-Driven CLP-Based test case generation. *Logic-Based Program Synthesis and Transformation* (2012)
17. J. Zhang, S.: Automated test case generation for the stress testing of multimedia systems. *Softw. Pract. Exper.* (2002)
18. L. C. Briand, Y. Labiche, M.: Using genetic algorithms for early schedulability analysis and stress testing in RT systems. *Genetic Programming and Evolvable Machines* (2006)
19. M. S. Bayan, J.: Automatic stress and load testing for embedded systems. *30th Annual International Computer Software and Applications Conference* (2006)
20. A. Avritzer, E.: Generating test suites for software load testing. *ACM SIGSOFT international symposium on Software testing and analysis* (1994)
21. A. Avritzer, E.: The automatic generation of load test suites and the assessment of the resulting software. *IEEE Trans. Softw. Eng.* (1995)
22. V. Garousi, L. C. Briand, Y.: Traffic-aware stress testing of distributed systems based on uml models. *International conference on Software engineering* (2006)
23. Jiang, Z.M., Hassan, A.E., Hamann, G., Flora, P.: Automated performance analysis of load tests. In: *IEEE International Conference on Software Maintenance*, Ieee (September 2009) 125–134
24. Haroon Malik, Bram Adams, A.: Pinpointing the subsys responsible for the performance deviations in a load test. *Software Reliability Engineering (ISSRE)* (2010)